

GraphML & TigerGraph

- 1) Query - Single Server Mode.
- * If your query starts from one or a few starting vertices, use this.

In this, the cluster elects one server to be master for that query.

All query computation takes place at query master.

Vertex and edge data are copied to the query master as needed.

- 2) Query - Distributed Mode

- * If your query starts from all or most vertices, use this mode.

In this, the server that received the query becomes the master.

Computation executes on all servers in parallel.

Accumulators are transferred across the cluster.

Page Rank:

A page has a high rank if the sum of the ranks of its backlinks is high.

(Backlinks: No. of web pages linking to webpage in focus).

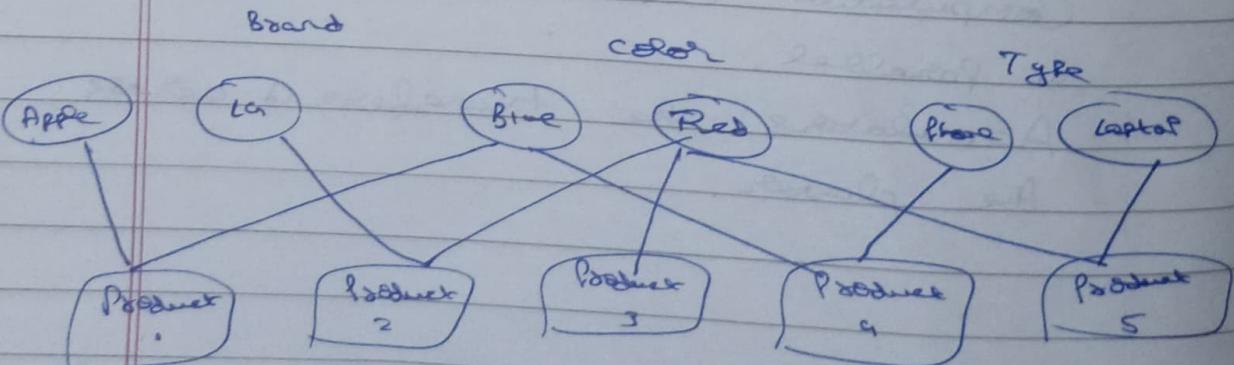
This covers both the case when a page has many backlinks and when a page has a few highly ranked backlinks.

Attribute or Vester?

It is beneficial to represent a column (attribute) as a vector type if you will frequently need to query for particular values of the property.

This way, the vertices act like an search index.

Ex: All the red products are connected to the Red vertex under Color type.



Accumulators

Accumulators are special type of variables that accumulate information about the graph during the traversal.

Phase 1

- (i) Receives messages, which will temporarily be put to a bucket that belongs to the accumulator.

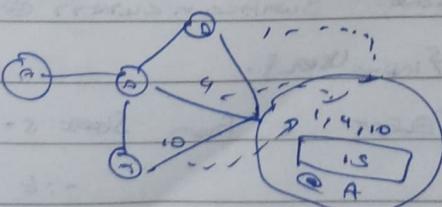
Phase 2

- (ii) The accumulator will aggregate the messages it received based on its accumulator type. The aggregated value will become the accumulators Value, and this value can be accessed.

1) Local Accumulators

- Each Selected vertex has its own accumulator.
- Local means per vertex. Each vertex does its own processing and considers what it can see / read / write.

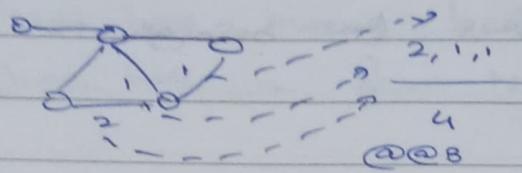
Ex: SumAccm@ A;



2) Global Accumulators

- Started in globally visible to all format.
- All vertices and edges have access.

Ex: SumAccum @@ B;



Some Accumulators

- 1) SumAccum <int>
- 2) MaxAccum <int>
- 3) MinAccum <int>
- 4) AvgAccum

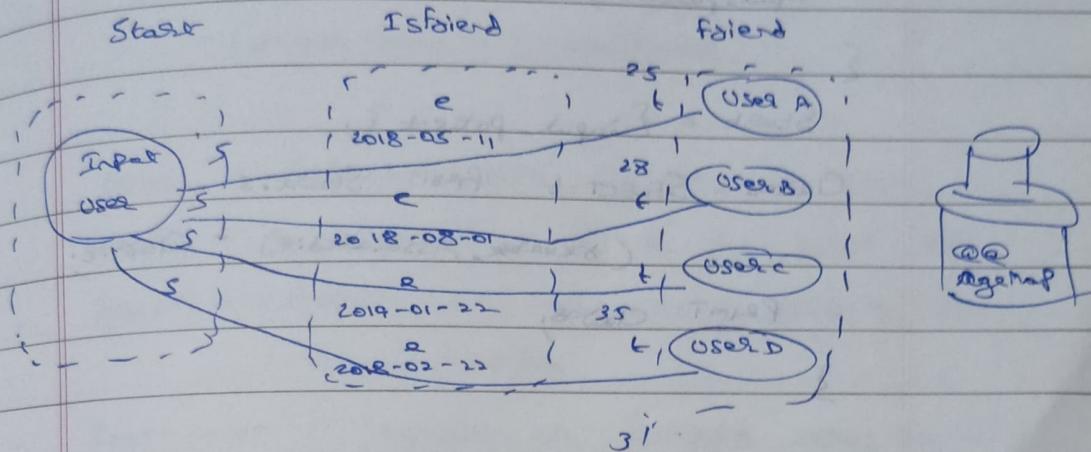
Ex. Accum Clause

Q) What is the age distribution of friends that were registered in 2018?

A) CREATE QUERY Getfriends (vertex <User> inputUser) FOR GRAPH Social {
 MapAccum <int>, SumAccum <int> @@ ageMap;
 Start = {inputUser};
 Friends = SELECT t FROM Start:s - (ISFriend:e)
 WHERE e.connectDt BETWEEN to_datetime("2018-01-01")
 AND to_datetime("2019-01-01")
 Accru @@ ageMap += (t.age/10->1);

PRINT @@ ageMap;

3



Note: Design the traversal plan to avoid skipping from the hub nodes.

→ By default 'As attribute' option is not enabled \Rightarrow ID value is not accessible from the query.

Therefore 'As attribute' needs to be checked when generic ID values are used in the query logic for non-joining purposes such as string concatenation, value comparison.

Having it switched off is more memory saving.

Cypher

Q) Find all the claims of a patient

A) CREATE QUERY `GetClaims` (`vertex = vseq`)
 input-patient) FOR GRAPH Social

{

`Start = {input-patient};`

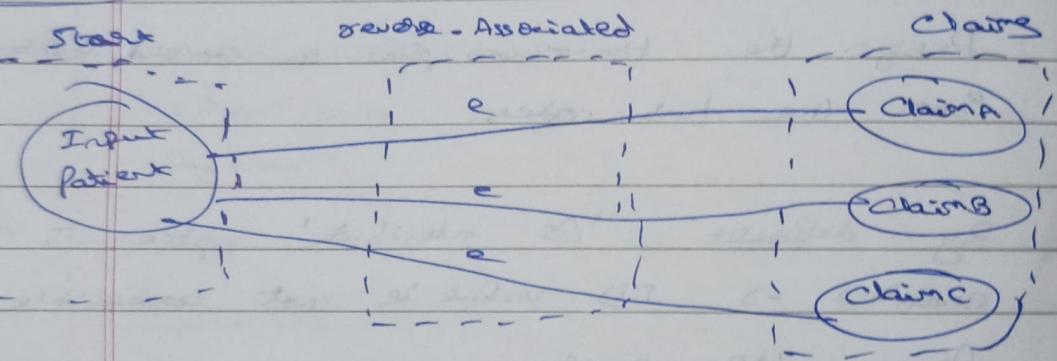
`Claims = SELECT t FROM Start:s -`

`(reverse-Associated:e) - Claim:t;`

`PRINT Claims;`

}

Visualisation



Explanation

- `Start` is a 'vertex set' initialized by the input vertex input-patient.
- `FROM` clause finds edges which match the pattern: source vertex is a Start, edge type is reverse-Associated, and target vertex is assigned to Claim type.
- `s, e, t` are aliases.
- `Claims` is a new vertex set equal to `t`.

*) result Set = $\text{SELECT } v \text{Set}$
 $\text{FROM } (\text{edgeSet} \mid \text{vertexSet})$
[where clause] [accum clause]
[postAccum clause] [having clause]
[order clause] {limit clause};

- FROM: Select active vertices & edges
 - WHERE: conditionally filter the active sets.
 - ACCUM: iterate on edge set; compute with accumulators
 - POST-ACCUM: iterate on vertex sets; compute with accumulators
 - HAVING: Conditionally filter the result set
 - ORDER BY: sort
 - LIMIT: max number of items
 - SELECT: result from source or target set
- What is within 'print' basis is returned as JSON.

result =
Ex: $\text{SELECT } s \text{ FROM } \text{Stations};$

Note: TigerGraph supports full ACID transactions with sequential consistency.

→ Main ways to interact with the Graph
Solutions are: REST endpoints, Graph Studio
and GSQL Shell.

Data Loading:

If you have defined a graph schema, you can load data into the graph store. Loading data is a 2-step process:

- i) Define a loading job, in which you use data loading statements to specify how values in the data source are extracted, transformed and loaded their destination.
- > ii) Run the loading job. Supply parameters if any.

Note:

GSQL code must be entered as a single line (unless the multi-line mode is used) within gsql shell

①

ls :

To check existing databases

②

DROP ALL :

Delete all database data, schema and all related definitions

Note: You can also run GSQL commands from a Linux shell by using 'gsql' followed by the 'command' enclosed in single quotes.

Ex: gsql 'ls'

The first step in creating a graph is to define its schema.

Cypher provides DDL commands for the same.

③

Creating a vertex type

Ex: CREATE VERTEX person (
 PRIMARY_ID name STRING,
 name STRING, age INT,
 gender STRING, state STRING
)

④

Create an Edge Type

It can be directed or undirected.

Ex:

CREATE UNDIRECTED EDGE friendship
 (FROM person, TO person, connect-day DATETIME
 : Property)

⑤

Creating a Graph

Here, we list the vertex types and edge types that we want to include in this graph.

Ex:

CREATE GRAPH social (person, friendship)
 :
 name of graph

⑥ SQl command to define a Loading Job

ER: USE GRAPH social;

BEGIN

CREATE LOADING JOB load-social

FOR GRAPH social {

DEFINE FILENAME file1 = "/home/tigergraph/person.csv";

DEFINE FILENAME file2 = "friendship.csv";

LOAD file1 TO VERTEX person VALUES

(\$ "name", \$ "name", \$ "age", \$ "gender",
\$ "state"), USING header = "true",
separator = ",";

for primary

for name
attribute

LOAD file2 TO EDGE friendship VALUES

(\$0, \$1, \$2) USING header = "true",
separator = ",";

}

END

Source column Position

⑦ Run the Loading Job

>>> RUN LOADING JOB load-social

→ Parameterized queries let you traverse the graph from one vertex set to an adjacent set of vertices, again and again, performing computations along the way, with built-in parallel execution and handy aggregation operations. You can even call one query from another.

Simple 1-Hop Queries

USE GRAPH Social
SET syntax-version = "v2"

CREATE QUERY hello (VERTEX persons p) {
 Stack = {p};
 Result = SELECT *
 FROM Stack - (Friendship: e)
 PRINT Result;

(/) Present will return the target vertex set

You can paste this query in a file with = gsql extension, say, hello.gsql

To join them, use >> @hello.gsoc2

Query is now in the Catalog.

Now, we have to install this:

>>> INSTALL QUERY hello

The installation will generate machine instructions and a REST endpoint and helps with faster query execution when executed multiple times.

Running the query:

>>> RUN QUERY hello ("Tom")

You can run queries without installing too (Anonymous queries)

Example 2

hello2.gsql

>>> USE GRAPH social

CREATE QUERY hello2 (VERTEX <person> p) {

Accum @visited = false;

Accum @@avgAge;

Start = {p};

FirstNeighbors = SELECT tgt

FROM Start:s - (friendship:e) -> person:tgt

Accum tgt:@visited += true,

s:@visited += true;

SecondNeighbors = SELECT tgt
 FROM FirstNeighbors - (:o) - :tgt
 WHERE tgt.@visited == false
 POST-ACCUM @@avg_Age += tgt.age;

PRINT SecondNeighbors;

PRINT @@avg_Age;

3

INSTALL QUERY hello2;

RUN QUERY hello2("Tom")

- The above query finds the average age of all neighbors 2-hops away from parameterized input person.
- To declare a local accumulator, we prefix an identifier name with a single "@" symbol.
- The identifier for a global accumulator begins with two "@"'s.
- The \pm operator within an ACCUM clause means that for each edge matching the FROM clause pattern, we accumulate the right-hand-side expression (true in this case) to the left-hand accumulator + (tgt.@visited and \$@visited).
- If edge type is omitted, like in (A), it is interpreted as ALL types.

* cURL is a command-line tool for transferring data using various network protocols.
It stands for Client URL.

Date	/ /
Page	bilt 10 on 10 Student Notebooks

→ Post_Accum traverses the vertex sets instead of the edge sets, guaranteeing that we do not double-count any vertices.

→ `rr SELECT count(*) FROM person`

→ `rrr SELECT count(*) FROM person - (Friendship) -> person`

→ `rrrr SELECT * FROM person WHERE primary_id = "Tom"`

→ `rrrr SELECT * FROM person - (Friendship) -> person
WHERE from_id == "Tom"`

(*) PATTERN MATCHING TUTORIAL

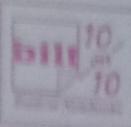
→ A graph pattern is a traversal base on the graph Schema. A pattern can be a linear trace or non-linear (tree, code, etc.)

Ex: Person - (Friendship) - Person - (Friendship) - Person

or, use +2 to denote two consecutive friendship edges,

Person - (Friendship +2) - Person

* Pattern matching is the process of finding subgraphs in a data graph that conform to a given query pattern.



Pattern matching queries supports nested querying.

Pattern example:

`Person:p = ((LIKES ?e HAS-CREATOR) :o)`

- Message: m

Indicates incoming 'HAS-CREATOR'

edge or outgoing 'LIKES' edge

or same querying expression

'|' is alternation separator

EDGES

→ '?' is a wildcard meaning any edge type.

→ The '()' empty pattern means any edge, directed or undirected.

VERTICES

* → Use '?', 'ANY' or () omit the vertex type to indicate any vertex type. If you omit the vertex type, you must provide a vertex alias.

Examples:

1) `SELECT t FROM Person:s - (IS-LOCATED-IN ?e) -City:t`

2) `SELECT t FROM (Post/Comment):s -`

`(IS-LOCATED-IN ?e) - Country:t`

3) `SELECT s FROM :s - (IS-LOCATED-IN ?e) -`
`Country:t`

4) `Select t FROM -(IS-LOCATED-IN ?e) -`
`Country:t`

Examples of 1-Hop patterns

1) From $X:x - (E2 \triangleright :e2) - Y:y$

Right directed edge e_2 binds to the source of $E2$; y binds to the target of $E2$.

2) From $X:x - (-:e) - Y:y$

Any undirected edge between a member of X and a member of Y .

3) From $X:x - (-\triangleright :e) - Y:y$

Any right directed edge with source in x and target in y .

4) From $X:x - ((\leftarrow 1-) :e) - Y:y$

Any left directed or any undirected

5) From $X:x - ((E1 \triangleright E2 \triangleright 1 \leftarrow E5) :e) - Y:y$

Any one of the three edge patterns

6) From $X:x - () - Y:y$

* Any edge (directed or undirected)

* Same as $(\leftarrow 1 \rightarrow 1 -)$

→ To find common qualities without interlinking:

INTERPART → QUERIES → FOR → GRAPH

graphstructure → SYNTAX vs { query, listing }

→ Increasing Querying Lineage Interlinked!

SET `query=INTERPART = 6000`

// with id = 1 PREVIOUSLY

② Find persons who know the person named "Victor" AND whose last name is "Black". Such persons.

③ USE GRAPH graphname

FROM "Persons" WHERE name = "Victor"

INTERPART → QUERY C) SYNTAX vs {

`friends = SELECT p`

`FROM persons = (knownas) = Person.p`

`WHERE s.first-name = "Victor" AND`

`s.last-name = "Black"`

`ORDER BY p.birthday ASC`

`LIMIT 3;`

PRINT friends[personid, first-name, friends.lastname,
friends.birthday];

3

Q) Find total number of comments and posts liked by viktor. A person can reach comments or posts via a directed edge LIKES.

A) USE GRAPH labc-sub

INTERPRET QUERY :: SYNTAX v2 {

SumAccum <@comment-cnt = 0;

SumAccum <@post-cnt = 0;

Result = SELECT s

FROM Person:s - (Likes?) - :tgt

WHERE s.first-name == "viktor" AND

s.lastname == "Akhiezer"

Accum CASE WHEN tgt.type == "Comment" THEN

s.@comment-cnt += 1

WHEN tgt.type == "Post" THEN

s.@post-cnt += 1

END;

PRINT Result [Result.@comment-cnt,

Result.@post-cnt],

}

→ From Person:tgt - (Has-Creator(Likes)) - (Comment|Post)

:src

Here, Has-Creator edge type starts from

Comment | Post, and AKES edge type

→ Starts from Person

→ Repeating 1-Hop Pattern

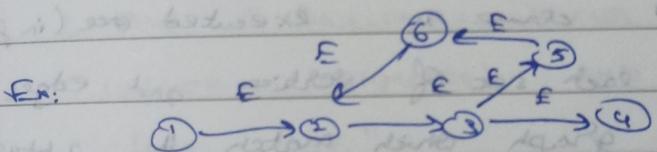
(i) FROM $x:x - (E \rightarrow \dots) - y:y$

→ Already chain of 2 E edges

(ii) FROM $x:x - (F \rightarrow \dots) - y:y$

→ Existing SF chain between 0 and 3 F edges

Note: Only the shortest matching occurrences are selected.



per 1 - (E > *) - q, b/a

$\rightarrow 1 \rightarrow 2 \rightarrow 3 \rightarrow 4 \quad \checkmark$

$\rightarrow 1 \rightarrow 2 \rightarrow 3 \rightarrow 5 \rightarrow 6 \rightarrow 2 \rightarrow 3 \rightarrow 4 \quad \times$

Q) "People who bought the product also bought the other product"

A) FROM This_Product:p - (< Bought:b1)

AND must Customer:c - (< Bought>:b2) - Product:p2

WHERE p2 = p

2-Hop pattern

FROM $x:n - (E1:e1) - Y:y - (E2:e2) - Z:z$

(or)

FROM $x:n - (E1,E2:z) - Z:z$

↓

Star intermediate vertex set is
when
not required. Vertex should not
be only of a specific type
" " means any vertex in b/w.

Note: If a pattern has a Kleene star (*)
to repeat an edge), SQL pattern
matching selects only the shortest
paths which match the pattern.

→ Accum clause is executed once (in parallel)
for each set of vertices and edges in
the graph which match the pattern
and constraints given in the FORM
and WHERE clauses.

→ Use Post-Accum clause to perform a
second, different kind of computation
on the results of your pattern matching.

Note:

→ The Accum clause executes for each
full path that matches the pattern in
the FROM clause.

In contrast, the Post-Accum clause
executes for each vertex in one vertex set.

→ , POST-ACCUM clause can use only 1 alias.

2) Find Nikita Arshavin's liked messages whose author's last names begin with S. find these author's overall count.

A) USE GRAPH lab-snb

INTERPRET QUERY () SYNTAX v2 { }

SumAccum <@> @@ok;

F = SELECT t

FROM :s - (Likes>:e) - :msg - (Has.Creator>)-

WHERE s.firstName == "Nikita" AND

s.lastName == "Arshavin" AND t.lastName

LIKE "%S%" ;

Answer = SELECT p

FROM Person - (Study-At>) - :u -

(< Study-At) - F:s

WHERE s != p AND s != u ()

Peer(p)

POST-ACCUM @@@ok + = 1;

PRINT @@@ok;

→ The PER clause specifies a list of vertex aliases, which are used to group the nodes in the match table, one group per distinct value of the alias or of the alias list.

→ Conjunctive Pattern Matching

Logical AND of 2 or more patterns that can be ~~be~~ naturally joined on a common alias.

Ex: SELECT FROM

Person: p - (KNOWS) - :tgt,

Posters - (<LIKES) - :tgt;

→ Data Modification

Pattern Matching GSQL supports Insert, Update and Delete operations.

Data modification happens in 2 ways using GSQL:

- 1) Top level. The statement does not need to be within any other statement.
- 2) Within a SELECT query statement. The FROM - ~~DELETE~~ WHERE clause defines a match table, and the data modification is performed based on the vertex and edge information in the match table.

The GSQL Specifications calls those
within-SQL statements DML-sub statements.

Data modification in interpreted mode
is not available yet.

SELECT queries with data modification
may only have one POST-Accum clause.

WSBPT

(1) Create a Person vertex, whose name is Tiger Woods. Next, find Victoria's favorite 2012 Poet authors, whose last name is prefixed with S. Finally, insert ~~Knows~~
~~Knows~~ edges connecting Tiger Woods with Victoria's favorite authors.

A) CREATE QUERY insert-edge-and-vertex ()
syntax v2 { selected var }

AA) INSERT INTO Person VALUES (10000000, "Tiger", "Woods",
"m", - , - , - , -)

R = SELECT t
FROM Person:s - (Likes >) - msg - (Has_Created >)
Person:f

WHERE s.firstname = "Victoria" AND s.lastname = "Akhilesh"
AND f.lastname like "S." AND year(
msg.creation_date) = 2012

PER (S, t)

Accum

BB)

INSERT INTO Knowe VALUES (100000000, t,
to_datetime('2020-06-01', ''))

PRINT r {r.id, r.first-name, r.last-name},
}

'AA' is called a top-level statement
'BB' is DML-sub INSERT statement.

UPDATE

Top-level UPDATE statements not supported yet.

Vertex attributes can only be updated
in Post-Accum clause, and edge
attributes can only be updated in
the Accum clause

To perform within-SELECT updates, the
FROM pattern can only be a
single hop, fixed length + pattern.

For all knows edges that connect user Ahmed and his friends whose last-name begins with "S", update the edge creation date to "2020-10-01".
 Also, for the person vertex (Tiger Woods) update the vertex's creation-date and language he speaks.

A) CREATE QUERY update_knows.ts (1) SYNTAX v2 {

R = SELECT p FROM Person:p

WHERE p.first-name == "Tiger" AND p.last-name == "Woods"

POST - ACCUM

// update simple base type attribute
 p.creation-date = to-datetime("2020-6-1")

// update collection-type attribute
 p.speaks = ("English", "Golf");

R = SELECT t

FROM Person:s-(knows:e)-:l

WHERE s.first-name == "Tiger" and
 s.last-name == "Woods"

Accum

e.creation-date =

to-datetime("2020-10-01");

3

DELETE

- Edges can only be deleted in the Accum clause.
- For best performance, vertices should be deleted in the Post-Accum clause.

Q) Delete vertex Tiger Woods and its known edges.

A) CREATE QUERY delete-edge-and-vertex()

SYNTAX :- {

R = SELECT t
 FROM Person:s - (Knows:e) - Person:t
 WHERE s.First-name = "Tiger" AND s.last-name = "Woods"
 Accum
 // delete edges
 DELETE (e)
 Post-Accum DELETE (s); // delete vertices

Print : R(R.id, R.first-name, R.last-name);
 3

→ For multi-line statements in "Shell",
 use the keywords BEGIN
 ...
 END.

CREATE VERTEX Syntax

`CREATE VERTEX Vertex-Type Name "("`
`primary_id -> name -> type`
`[, " attribute_name type (DEFAULT default_value)] + ")"`

`[WITH [STATUS = "none" | "outdegree_by_edgetype"]`

`[primary_id-as attribute = "true"]]`

Ex: `CREATE VERTEX Movie (PRIMARY_ID id INT,`
`name STRING, year INT)`

`WITH primary_id-as attribute = "true"`

→ GSQL creates a hash index on primary id
 with $O(1)$ time complexity.

`CREATE VERTEX Movie (id INT PRIMARY KEY,`
`name STRING, year INT)`

Composite Keys

→ `CREATE VERTEX Movie (id INT, title`
`STRING, year INT, PRIMARY KEY (title, year, id))`

Creating Edges:

`CREATE DIRECTED EDGE Member_OF`
`(FROM Person, TO Dog | FROM Dog, TO Dog,`
`joined DATETIME) ... [WITH REVERSE_EDGE`
`as "Rev_Many"]`

It can be UNDIRECTED edge

→ Checking a graph with all vertex & edge types

Ex: CREATE GRAPH graph-name (+)

→ altering vertex attribute:

Ex: ALTER VERTEX Company ADD ATTRIBUTE
(industry STRING, market-cap Double)

For above,

you have to create Schema job
and run the schema job

*

Centrality Algorithms:

Centrality algorithms find the important vertices in a graph based on their connections with other vertices.

A centrality algorithm usually assigns

* scores to vertices or edges based on how close they are to the center of connection-based activity.

The definition of closeness and what constitutes the center of connection-based activity depend on the algorithm.

Ex: Degree centrality

* Classification Algorithms

They assign a label (class name) to a vertex based on it satisfying the established conditions for membership in that class.

This differs from ^{*} community (or clustering) algorithms in which there are no pre-established conditions for membership.

* Node Embedding Algorithms

Node embeddings are vector representations of properties of vertices in a graph. These vectors can be used for machine learning.

* Pathfinding Algorithms

It finds the best path(s) between two vertices or among a set of vertices according to a particular rule.

* Similarity Algorithms

Similarity Algorithms assign a score to a pair of vertices based on how similar they are. This usually involves seeing if they have the same or similar neighborhoods. This section also includes distance algorithm, as distance is the inverse of similarity.

#) Topological Walk Prediction Algorithm
These algorithms determine the closeness score of a pair of vertices. This score can then be used to make predictions about the relationship between the two vertices in the pair.

Ex: Adamic Adar, Common Neighbors